

Welcome and Introduction

Philip Schulz and Wilker Aziz

<https://github.com/philschulz/VITutorial>

About us . . .

Wilker Aziz

- ▶ Assistant Professor at UvA-ILLC
- ▶ VI, sampling methods, NLP

Philip Schulz

- ▶ Applied Scientist at Amazon
- ▶ VI, machine translation, Bayesian models

Why are we here today?

Because NN models work

Why are we here today?

Because NN models work
but they may struggle with

- ▶ lack of supervision
- ▶ partial supervision
- ▶ lack of inductive bias

Why are we here today?

Because NN models work
but they may struggle with

- ▶ lack of supervision
- ▶ partial supervision
- ▶ lack of inductive bias

In such cases, we usually turn to probabilistic models

Why are we here today?

Because NN models work
but they may struggle with

- ▶ lack of supervision
- ▶ partial supervision
- ▶ lack of inductive bias

In such cases, we usually turn to probabilistic models
but it turns out NNs and PGMs do not go along

- ▶ well...

Why are we here today?

Because NN models work
but they may struggle with

- ▶ lack of supervision
- ▶ partial supervision
- ▶ lack of inductive bias

In such cases, we usually turn to probabilistic models
but it turns out NNs and PGMs do not go along

- ▶ well... that's no longer true!

What are you getting out of this today?

As we progress we will

- ▶ develop a shared vocabulary to talk about generative models powered by NNs
- ▶ derive crucial results step by step

What are you getting out of this today?

As we progress we will

- ▶ develop a shared vocabulary to talk about generative models powered by NNs
- ▶ derive crucial results step by step

Goal

- ▶ you should be able to navigate through fresh literature
- ▶ and start combining probabilistic models and NNs

Supervised problems

We have data $x^{(1)}, \dots, x^{(N)}$ e.g. sentences, images generated by some **unknown** procedure which we assume can be captured by a probabilistic model

- ▶ with **known** probability (mass/density) function e.g.

$$X \sim \text{Cat}(\theta_1, \dots, \theta_K) \quad \text{or} \quad X \sim \mathcal{N}(\theta_\mu, \theta_\sigma^2)$$

Supervised problems

We have data $x^{(1)}, \dots, x^{(N)}$ e.g. sentences, images generated by some **unknown** procedure which we assume can be captured by a probabilistic model

- ▶ with **known** probability (mass/density) function e.g.

$$X \sim \text{Cat}(\theta_1, \dots, \theta_K) \quad \text{or} \quad X \sim \mathcal{N}(\theta_\mu, \theta_\sigma^2)$$

and **estimate parameters** θ that assign maximum likelihood $p(x^{(1)}, \dots, x^{(N)} | \theta)$ to observations

Supervised NN models

Let y be all side information available
e.g. deterministic *inputs/features/predictors*

Have neural networks predict parameters of our probabilistic model

$$X|y \sim \text{Cat}(\pi_{\theta}(y)) \quad \text{or} \quad X|y \sim \mathcal{N}(\mu_{\theta}(y), \sigma_{\theta}(y)^2)$$

and proceed to **estimate parameters θ** of the NNs


Graphical model

Random variables


- ▶ observed data
 $x^{(1)}, \dots, x^{(N)}$

Deterministic variables

- ▶ inputs or predictors
 $y^{(1)}, \dots, y^{(N)}$
- ▶ model parameters θ



Multiple problems, same language



(Conditional) Density estimation

Parsing	Side information (y) a sentence	Observation (x) its syntactic/semantic parse tree/graph
Translation	a sentence	its translation
Captioning	an image	caption in English
Entailment	a text and hypothesis	entailment relation

Task-driven feature extraction

Often our side information is itself some high dimensional data

- ▶ y is a sentence and x a tree
- ▶ y is the source sentence and x is the target
- ▶ y is an image and x is a caption

and part of the job of the NNs that parametrise our models is to also **deterministically** encode that input in a low-dimensional space

NN as efficient parametrisation

From a statistical point of view, NNs do not generate data

- ▶ they parametrise distributions that *by assumption* govern data
- ▶ compact and efficient way to map from complex side information to parameter space

NN as efficient parametrisation

From a statistical point of view, NNs do not generate data

- ▶ they parametrise distributions that *by assumption* govern data
- ▶ compact and efficient way to map from complex side information to parameter space

Prediction is done by a decision rule outside the statistical model

- ▶ e.g. argmax, beam search

Maximum likelihood estimation

Let $p(x|\theta)$ be the probability of an observation x
and θ refer to all of its parameters

Maximum likelihood estimation

Let $p(x|\theta)$ be the probability of an observation x
and θ refer to all of its parameters

Given a dataset $x^{(1)}, \dots, x^{(N)}$ of i.i.d. observations,

Maximum likelihood estimation

Let $p(x|\theta)$ be the probability of an observation x and θ refer to all of its parameters

Given a dataset $x^{(1)}, \dots, x^{(N)}$ of i.i.d. observations, the log-likelihood function gives us a criterion for parameter estimation

$$\mathcal{L}(\theta|x^{(1:N)}) =$$

Maximum likelihood estimation

Let $p(x|\theta)$ be the probability of an observation x and θ refer to all of its parameters

Given a dataset $x^{(1)}, \dots, x^{(N)}$ of i.i.d. observations, the log-likelihood function gives us a criterion for parameter estimation

$$\mathcal{L}(\theta|x^{(1:N)}) = \log \prod_{s=1}^N p(x^{(s)}|\theta) =$$

Maximum likelihood estimation

Let $p(x|\theta)$ be the probability of an observation x and θ refer to all of its parameters

Given a dataset $x^{(1)}, \dots, x^{(N)}$ of i.i.d. observations, the log-likelihood function gives us a criterion for parameter estimation

$$\mathcal{L}(\theta|x^{(1:N)}) = \log \prod_{s=1}^N p(x^{(s)}|\theta) = \sum_{s=1}^N \log p(x^{(s)}|\theta)$$

MLE via gradient-based optimisation

If the log-likelihood is **differentiable** and **tractable** then backpropagation gives us the gradient

$$\nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) =$$

MLE via gradient-based optimisation

If the log-likelihood is **differentiable** and **tractable** then backpropagation gives us the gradient

$$\nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) = \nabla_{\theta} \sum_{s=1}^N \log p(x^{(s)} | \theta) =$$

MLE via gradient-based optimisation

If the log-likelihood is **differentiable** and **tractable** then backpropagation gives us the gradient

$$\nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) = \nabla_{\theta} \sum_{s=1}^N \log p(x^{(s)} | \theta) = \sum_{s=1}^N \nabla_{\theta} \log p(x^{(s)} | \theta)$$

MLE via gradient-based optimisation

If the log-likelihood is **differentiable** and **tractable** then backpropagation gives us the gradient

$$\nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) = \nabla_{\theta} \sum_{s=1}^N \log p(x^{(s)} | \theta) = \sum_{s=1}^N \nabla_{\theta} \log p(x^{(s)} | \theta)$$

and we can update θ in the direction

$$\gamma \nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)})$$

to attain a local maximum of the likelihood function

Big Data

For large N , computing the gradient is inconvenient

$$\nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) = \underbrace{\sum_{s=1}^N \nabla_{\theta} \log p(x^{(s)} | \theta)}_{\text{too many terms}}$$

Big Data

For large N , computing the gradient is inconvenient

$$\begin{aligned}\nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) &= \underbrace{\sum_{s=1}^N \nabla_{\theta} \log p(x^{(s)} | \theta)}_{\text{too many terms}} \\ &= \sum_{s=1}^N \frac{1}{N} N \nabla_{\theta} \log p(x^{(s)} | \theta)\end{aligned}$$

Big Data

For large N , computing the gradient is inconvenient

$$\begin{aligned}
 \nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) &= \underbrace{\sum_{s=1}^N \nabla_{\theta} \log p(x^{(s)} | \theta)}_{\text{too many terms}} \\
 &= \sum_{s=1}^N \frac{1}{N} N \nabla_{\theta} \log p(x^{(s)} | \theta) \\
 &= \sum_{s=1}^N \mathcal{U}(s|1/N) N \nabla_{\theta} \log p(x^{(s)} | \theta)
 \end{aligned}$$

Big Data

For large N , computing the gradient is inconvenient

$$\begin{aligned}
 \nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) &= \underbrace{\sum_{s=1}^N \nabla_{\theta} \log p(x^{(s)} | \theta)}_{\text{too many terms}} \\
 &= \sum_{s=1}^N \frac{1}{N} N \nabla_{\theta} \log p(x^{(s)} | \theta) \\
 &= \sum_{s=1}^N \mathcal{U}(s | 1/N) N \nabla_{\theta} \log p(x^{(s)} | \theta) \\
 &= \mathbb{E}_{S \sim \mathcal{U}(1/N)} [N \nabla_{\theta} \log p(x^{(S)} | \theta)]
 \end{aligned}$$

S selects data points uniformly at random

Stochastic optimisation

For large N , we can use a gradient estimate

$$\nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) = \underbrace{\mathbb{E}_{S \sim \mathcal{U}(1/N)} [N \nabla_{\theta} \log p(x^{(S)} | \theta)]}_{\text{expected gradient :)}}$$

Stochastic optimisation

For large N , we can use a gradient estimate

$$\nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) = \underbrace{\mathbb{E}_{S \sim \mathcal{U}(1/N)} [N \nabla_{\theta} \log p(x^{(S)} | \theta)]}_{\text{expected gradient :)}}$$

$$\stackrel{\text{MC}}{\approx} \frac{1}{M} \sum_{m=1}^M N \nabla_{\theta} \log p(x^{(s_m)} | \theta)$$

$$S_m \sim \mathcal{U}(1/N)$$

Stochastic optimisation

For large N , we can use a gradient estimate

$$\nabla_{\theta} \mathcal{L}(\theta | x^{(1:N)}) = \underbrace{\mathbb{E}_{S \sim \mathcal{U}(1/N)} [N \nabla_{\theta} \log p(x^{(S)} | \theta)]}_{\text{expected gradient :)}}$$

$$\stackrel{\text{MC}}{\approx} \frac{1}{M} \sum_{m=1}^M N \nabla_{\theta} \log p(x^{(s_m)} | \theta)$$

$$S_m \sim \mathcal{U}(1/N)$$

and take a step in the direction

$$\gamma \frac{N}{M} \underbrace{\nabla_{\theta} \mathcal{L}(\theta | x^{(s_1:s_M)})}_{\text{stochastic gradient}}$$

where $x^{(s_1:s_M)}$ is a random mini-batch of size M

DL in NLP recipe

Maximum likelihood estimation

- ▶ tells you which **loss** to optimise (i.e. negative log-likelihood)

Automatic differentiation (*backprop*)

- ▶ “give me a tractable forward pass and I will give you **gradients**”

Stochastic optimisation powered by backprop

- ▶ general purpose gradient-based optimisers

Constraints

Differentiability

- ▶ intermediate representations must be continuous
- ▶ activations must be differentiable

Tractability

- ▶ the likelihood function must be evaluated exactly, thus it's required to be tractable

When do we have intractable likelihood?

Unsupervised problems contain unobserved random variables

$$p(x, z|\theta) = \overbrace{p(z)}^{\text{prior}} \underbrace{p(x|z, \theta)}_{\text{observation model}}$$

When do we have intractable likelihood?

Unsupervised problems contain unobserved random variables

$$p(x, z|\theta) = \overbrace{p(z)}^{\text{prior}} \underbrace{p(x|z, \theta)}_{\text{observation model}}$$


thus assessing the marginal likelihood requires
marginalisation of latent variables

$$p(x|\theta) = \int p(x, z|\theta) \, dz = \int p(z)p(x|z, \theta) \, dz$$

Latent variable model

Latent random variables


- ▶ unobserved
- ▶ or unobservable



Latent variable model

Latent random variables

- ▶ unobserved
- ▶ or unobservable



A joint distribution over data and unknowns

$$p(x, z|\theta) = p(z)p(x|z, \theta)$$

Examples of latent variable models

Discrete latent variable, continuous observation

$$p(x|\theta) = \underbrace{\sum_{c=1}^K \text{Cat}(c|\pi_1, \dots, \pi_K)}_{\text{too many forward passes}} \underbrace{\mathcal{N}(x|\mu_\theta(c), \sigma_\theta(c)^2)}_{\text{forward pass}}$$

Examples of latent variable models

Discrete latent variable, continuous observation

$$p(x|\theta) = \underbrace{\sum_{c=1}^K \text{Cat}(c|\pi_1, \dots, \pi_K)}_{\text{too many forward passes}} \underbrace{\mathcal{N}(x|\mu_\theta(c), \sigma_\theta(c)^2)}_{\text{forward pass}}$$

Continuous latent variable, discrete observation

$$p(x|\theta) = \underbrace{\int \mathcal{N}(z|0, I)}_{\text{infinitely many forward passes}} \underbrace{\text{Cat}(x|\pi_\theta(z))}_{\text{forward pass}} dz$$

Intractable gradient

$$\nabla_{\theta} \log p(x|\theta)$$

Intractable gradient

$$\nabla_{\theta} \log p(x|\theta) = \nabla_{\theta} \log \underbrace{\int p(x, z|\theta) \, dz}_{\text{marginal}}$$

Intractable gradient

$$\begin{aligned}\nabla_{\theta} \log p(x|\theta) &= \nabla_{\theta} \log \underbrace{\int p(x, z|\theta) \, dz}_{\text{marginal}} \\ &= \underbrace{\frac{1}{\int p(x, z|\theta) \, dz} \int \nabla_{\theta} p(x, z|\theta) \, dz}_{\text{chain rule}}\end{aligned}$$

Intractable gradient

$$\begin{aligned}
 \nabla_{\theta} \log p(x|\theta) &= \nabla_{\theta} \log \underbrace{\int p(x, z|\theta) \, dz}_{\text{marginal}} \\
 &= \underbrace{\frac{1}{\int p(x, z|\theta) \, dz} \int \nabla_{\theta} p(x, z|\theta) \, dz}_{\text{chain rule}} \\
 &= \frac{1}{p(x|\theta)} \int \underbrace{p(x, z|\theta) \nabla_{\theta} \log p(x, z|\theta)}_{\text{log-identity for derivatives}} \, dz
 \end{aligned}$$

Intractable gradient

$$\begin{aligned}
 \nabla_{\theta} \log p(x|\theta) &= \nabla_{\theta} \log \underbrace{\int p(x, z|\theta) \, dz}_{\text{marginal}} \\
 &= \frac{1}{\underbrace{\int p(x, z|\theta) \, dz}_{\text{chain rule}}} \underbrace{\int \nabla_{\theta} p(x, z|\theta) \, dz}_{\text{chain rule}} \\
 &= \frac{1}{p(x|\theta)} \int \underbrace{p(x, z|\theta) \nabla_{\theta} \log p(x, z|\theta)}_{\text{log-identity for derivatives}} \, dz \\
 &= \int p(z|x, \theta) \nabla_{\theta} \log p(x, z|\theta) \, dz
 \end{aligned}$$

Gradient estimates?

$$\nabla_{\theta} \log p(x|\theta) = \int p(z|x, \theta) \nabla_{\theta} \log p(x, z|\theta) \, dz$$

Gradient estimates?

$$\begin{aligned}\nabla_{\theta} \log p(x|\theta) &= \int p(z|x, \theta) \nabla_{\theta} \log p(x, z|\theta) \, dz \\ &= \mathbb{E}_{p(z|x, \theta)} [\nabla_{\theta} \log p(x, z|\theta)]\end{aligned}$$

Gradient estimates?

$$\begin{aligned}\nabla_{\theta} \log p(x|\theta) &= \int p(z|x, \theta) \nabla_{\theta} \log p(x, z|\theta) \, dz \\ &= \mathbb{E}_{p(z|x, \theta)} [\nabla_{\theta} \log p(x, z|\theta)] \\ &\stackrel{\text{MC}}{\approx} \frac{1}{K} \sum_{k=1}^K \nabla_{\theta} \log p(x, z_k|\theta) \quad \text{where } z_k \sim p(z|x, \theta)\end{aligned}$$

Gradient estimates?

$$\begin{aligned}
 \nabla_{\theta} \log p(x|\theta) &= \int p(z|x, \theta) \nabla_{\theta} \log p(x, z|\theta) \, dz \\
 &= \mathbb{E}_{p(z|x, \theta)} [\nabla_{\theta} \log p(x, z|\theta)] \\
 &\stackrel{\text{MC}}{\approx} \frac{1}{K} \sum_{k=1}^K \nabla_{\theta} \log p(x, z_k|\theta) \quad \text{where } z_k \sim p(z|x, \theta)
 \end{aligned}$$

But what are we sampling from exactly?

$$p(z|x, \theta) =$$

Gradient estimates?

$$\begin{aligned}
 \nabla_{\theta} \log p(x|\theta) &= \int p(z|x, \theta) \nabla_{\theta} \log p(x, z|\theta) \, dz \\
 &= \mathbb{E}_{p(z|x, \theta)} [\nabla_{\theta} \log p(x, z|\theta)] \\
 &\stackrel{\text{MC}}{\approx} \frac{1}{K} \sum_{k=1}^K \nabla_{\theta} \log p(x, z_k|\theta) \quad \text{where } z_k \sim p(z|x, \theta)
 \end{aligned}$$

But what are we sampling from exactly?

$$p(z|x, \theta) = \frac{p(x, z|\theta)}{p(x|\theta)}$$

But why latent variable modelling?

Some reasons

- ▶ better handle on statistical assumptions
e.g. breaking marginal independence

But why latent variable modelling?

Some reasons

- ▶ better handle on statistical assumptions
e.g. breaking marginal independence
- ▶ organise a massive collection of data
e.g. LDA

But why latent variable modelling?

Some reasons

- ▶ better handle on statistical assumptions
e.g. breaking marginal independence
- ▶ organise a massive collection of data
e.g. LDA
- ▶ learn from unlabelled data
e.g. semi-supervised learning

But why latent variable modelling?

Some reasons

- ▶ better handle on statistical assumptions
e.g. breaking marginal independence
- ▶ organise a massive collection of data
e.g. LDA
- ▶ learn from unlabelled data
e.g. semi-supervised learning
- ▶ induce discrete representations
e.g. parse trees, dependency graphs, alignments

But why latent variable modelling?

Some reasons

- ▶ better handle on statistical assumptions
e.g. breaking marginal independence
- ▶ organise a massive collection of data
e.g. LDA
- ▶ learn from unlabelled data
e.g. semi-supervised learning
- ▶ induce discrete representations
e.g. parse trees, dependency graphs, alignments
- ▶ uncertainty quantification
e.g. Bayesian NNs

Examples: Lexical alignment

Generate a word x_i in L1 from a word y_{a_i} in L2

$$P(x|y, \theta) \stackrel{\text{ind}}{=} \prod_{i=1}^{|x|} \sum_{a_i=1}^{|y|} P(a_i|y) P(x_i|y_{a_i})$$

a mixture model whose mixture components are
labelled by words marginalisation $O(|x||y|)$

Examples: Rationale extraction

Sentiment analysis based on a subset of the input

$$P(x|y, \theta) = \sum_{f_1=0}^1 \cdots \sum_{f_{|y|=0}}^1 P(x|f, y) \prod_{i=1}^{|y|} \text{Bernoulli}(f_i|\theta_{y_i})$$

where $P(x|f, y)$ conditions on y_i iff $f_i = 1$.

A factor model whose factors are labelled by words
 marginalisation $O(2^{|y|})$

Examples: Language modelling

A (deterministic) RNNLM always produces the same conditional for a given prefix. Isn't it reasonable to expect the conditional to depend on what we are talking about?: e.g. *Rio de Janeiro ...*

- ▶ history: *was once the Brazilian capital*
- ▶ tourism: *offers some of Brazil's most iconic landscapes*
- ▶ news: *recently hosted the world cup final*

$$P(x|\theta) = \int \mathcal{N}(z|0, I) \prod_{i=1}^{|x|} P(x_i|z, x_{<i}, \theta)$$

Deep Generative Models

Probabilistic models parametrised by neural networks

Deep Generative Models

Probabilistic models parametrised by neural networks

- ▶ explicit modelling assumptions
one of the reasons why there's so much interest

Deep Generative Models

Probabilistic models parametrised by neural networks

- ▶ explicit modelling assumptions
one of the reasons why there's so much interest
- ▶ but requires efficient inference

Deep Generative Models

Probabilistic models parametrised by neural networks

- ▶ explicit modelling assumptions
one of the reasons why there's so much interest
- ▶ but requires efficient inference
which is the reason why we are here today