

Deep Generative Models: Discrete Latente Variables

Philip Schulz and Wilker Aziz

[https:
//github.com/philschulz/VITutorial](https://github.com/philschulz/VITutorial)

First Attempt: Wake-Sleep

Neural Variational Inference

Score function estimator

Variance reduction

Generative Models

Joint distribution over observed data x and latent variables Z .

$$p(x, z|\theta) = \underbrace{p(z)}_{\text{prior}} \underbrace{p(x|z, \theta)}_{\text{likelihood}}$$

The likelihood and prior are often standard distributions (Gaussian, Bernoulli) with simple dependence on conditioning information.

Deep generative models

Joint distribution with **deep observation model**

$$p(x, z|\theta) = \underbrace{p(z)}_{\text{prior}} \underbrace{p(x|z, \theta)}_{\text{likelihood}}$$

mapping from z to $p(x|z, \theta)$ is a NN with parameters θ

Deep generative models

Joint distribution with **deep observation model**

$$p(x, z|\theta) = \underbrace{p(z)}_{\text{prior}} \underbrace{p(x|z, \theta)}_{\text{likelihood}}$$

mapping from z to $p(x|z, \theta)$ is a NN with parameters θ

Marginal likelihood

$$p(x|\theta) = \int p(x, z|\theta) \, dz = \int p(z)p(x|z, \theta) \, dz$$

intractable in general

Goals

We want

- ▶ richer probabilistic models

Goals

We want

- ▶ richer probabilistic models
- ▶ complex observation models
parameterised by NNs

Goals

We want

- ▶ richer probabilistic models
- ▶ complex observation models
parameterised by NNs

but we can't perform gradient-based MLE

Goals

We want

- ▶ richer probabilistic models
- ▶ complex observation models
parameterised by NNs

but we can't perform gradient-based MLE

We need **approximate inference** techniques!

First Attempt: Wake-Sleep

Neural Variational Inference

Score function estimator

Variance reduction

Wake-sleep Algorithm

- ▶ Generalise latent variables to Neural Networks
- ▶ Train generative neural model
- ▶ Use variational inference! (kind of)

Wake-sleep Architecture

2 Neural Networks:

Wake-sleep Architecture

2 Neural Networks:

- ▶ A generation network to model the data (the one we want to optimise) – parameters: θ

Wake-sleep Architecture

2 Neural Networks:

- ▶ A generation network to model the data (the one we want to optimise) – parameters: θ
- ▶ An inference (recognition) network (to model the latent variable) – parameters: λ

Wake-sleep Architecture

2 Neural Networks:

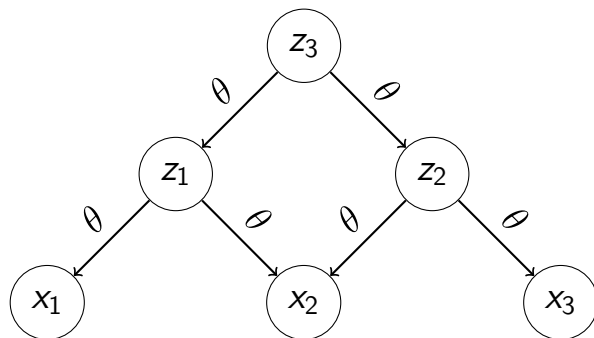
- ▶ A generation network to model the data (the one we want to optimise) – parameters: θ
- ▶ An inference (recognition) network (to model the latent variable) – parameters: λ
- ▶ Original setting: binary hidden units

Wake-sleep Architecture

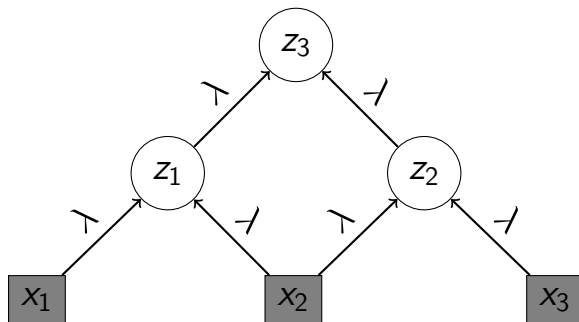
2 Neural Networks:

- ▶ A generation network to model the data (the one we want to optimise) – parameters: θ
- ▶ An inference (recognition) network (to model the latent variable) – parameters: λ
- ▶ Original setting: binary hidden units
- ▶ Training is performed in a “hard EM” fashion

Generator



Inference Network



Wake-sleep Training

Wake Phase

- ▶ Use inference network to sample hidden unit setting z from $q(z|x, \lambda)$
- ▶ Update generation parameters θ to maximize joint log-likelihood of data and latents $p(x, z|\theta)$

Wake-sleep Training

Wake Phase

- ▶ Use inference network to sample hidden unit setting z from $q(z|x, \lambda)$
- ▶ Update generation parameters θ to maximize joint log-likelihood of data and latents $p(x, z|\theta)$

Sleep Phase

- ▶ Produce dream sample \tilde{x} from random hidden unit z
- ▶ Update inference parameters λ to maximize probability of latent state $q(z|\tilde{x}, \lambda)$

Wake Phase Objective

Objective

$$\arg \min_{\theta} \text{KL} (q(z|x, \lambda) || p(z|x, \theta))$$

Wake Phase Objective

Objective

$$\arg \min_{\theta} \text{KL} (q(z|x, \lambda) \parallel p(z|x, \theta))$$

$$= \arg \max_{\theta} \underbrace{\mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \mathbb{H}[q(z|x, \lambda)]}_{\mathcal{G}(\theta)}$$

Wake Phase Objective

Objective

$$\arg \min_{\theta} \text{KL} (q(z|x, \lambda) \parallel p(z|x, \theta))$$

$$= \arg \max_{\theta} \underbrace{\mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \mathbb{H}[q(z|x, \lambda)]}_{\mathcal{G}(\theta)}$$

Gradient estimate

$$\nabla_{\theta} \mathcal{G}(\theta) = \nabla_{\theta} \mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \nabla_{\theta} \mathbb{H}[q(z|x, \lambda)]$$

Wake Phase Objective

Objective

$$\begin{aligned} & \arg \min_{\theta} \text{KL} (q(z|x, \lambda) \parallel p(z|x, \theta)) \\ &= \arg \max_{\theta} \underbrace{\mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \mathbb{H}[q(z|x, \lambda)]}_{\mathcal{G}(\theta)} \end{aligned}$$

Gradient estimate

$$\begin{aligned} \nabla_{\theta} \mathcal{G}(\theta) &= \nabla_{\theta} \mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \nabla_{\theta} \mathbb{H}[q(z|x, \lambda)] \\ &= \mathbb{E}_{q(z|x, \lambda)} [\nabla_{\theta} \log p(z, x|\theta)] \end{aligned}$$

Wake Phase Objective

Objective

$$\begin{aligned} & \arg \min_{\theta} \text{KL} (q(z|x, \lambda) \parallel p(z|x, \theta)) \\ &= \arg \max_{\theta} \underbrace{\mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \mathbb{H}[q(z|x, \lambda)]}_{\mathcal{G}(\theta)} \end{aligned}$$

Gradient estimate

$$\begin{aligned} \nabla_{\theta} \mathcal{G}(\theta) &= \nabla_{\theta} \mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \nabla_{\theta} \mathbb{H}[q(z|x, \lambda)] \\ &= \mathbb{E}_{q(z|x, \lambda)} [\nabla_{\theta} \log p(z, x|\theta)] \\ &\stackrel{\text{MC}}{\approx} \nabla_{\theta} \log p(z, x|\theta) \quad \text{where } z \sim q(z|x, \lambda) \end{aligned}$$

Wake Phase Objective


Assumes latent state z to be fixed random draws from $q(z|x, \lambda)$.

$$\arg \min_{\theta} \text{KL} (q(z|x, \lambda) || p(z|x, \theta))$$
$$\stackrel{\text{MC}}{\approx} \arg \max_{\theta} \log p(z, x|\theta)$$

This is simply supervised learning with imputed latent data!

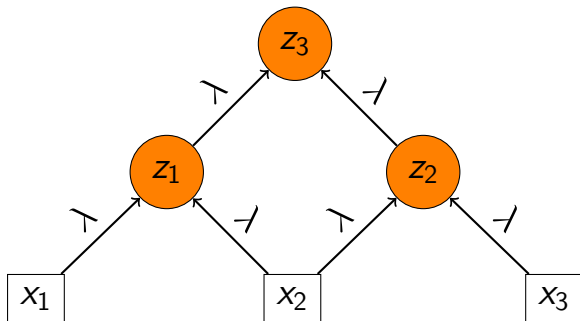
Wake Phase Sampling

Sampling $z \sim q(z|x, \lambda)$




Wake Phase Sampling

Sampling $z \sim q(z|x, \lambda)$



Wake Phase Update

Update θ



Sleep Phase Objective

Objective

$$\arg \min_{\lambda} \text{KL} (q(z|x, \lambda) || p(z|x, \theta))$$

Sleep Phase Objective

Objective

$$\begin{aligned} & \arg \min_{\lambda} \text{KL} (q(z|x, \lambda) || p(z|x, \theta)) \\ &= \arg \max_{\lambda} \underbrace{\mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \mathbb{H}[q(z|x, \lambda)]}_{\mathcal{R}(\lambda)} \end{aligned}$$

Sleep Phase Objective

Objective

$$\begin{aligned} & \arg \min_{\lambda} \text{KL} (q(z|x, \lambda) \parallel p(z|x, \theta)) \\ &= \arg \max_{\lambda} \underbrace{\mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \mathbb{H}[q(z|x, \lambda)]}_{\mathcal{R}(\lambda)} \end{aligned}$$

Gradient estimate

$$\nabla_{\lambda} \mathcal{R}(\lambda) = \nabla_{\lambda} \mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \nabla_{\lambda} \mathbb{H}[q(z|x, \lambda)]$$

Sleep Phase Objective

Objective

$$\begin{aligned} & \arg \min_{\lambda} \text{KL} (q(z|x, \lambda) \parallel p(z|x, \theta)) \\ &= \arg \max_{\lambda} \underbrace{\mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \mathbb{H}[q(z|x, \lambda)]}_{\mathcal{R}(\lambda)} \end{aligned}$$

Gradient estimate

$$\nabla_{\lambda} \mathcal{R}(\lambda) = \nabla_{\lambda} \mathbb{E}_{q(z|x, \lambda)} [\log p(z, x|\theta)] + \nabla_{\lambda} \mathbb{H}[q(z|x, \lambda)]$$

Let's change the objective!

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\arg \min_{\lambda} \mathbb{E}_{p(x)} [\text{KL} (p(z|x, \theta) \parallel q(z|x, \lambda))]$$

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\begin{aligned} & \arg \min_{\lambda} \mathbb{E}_{p(x)} [\text{KL} (p(z|x, \theta) \parallel q(z|x, \lambda))] \\ &= \arg \min_{\lambda} \mathbb{E}_{p(x, z|\theta)} [\log p(z|x, \theta) - \log q(z|x, \lambda)] \end{aligned}$$

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\begin{aligned}
 & \arg \min_{\lambda} \mathbb{E}_{p(x)} [\text{KL} (p(z|x, \theta) \parallel q(z|x, \lambda))] \\
 &= \arg \min_{\lambda} \mathbb{E}_{p(x, z|\theta)} [\log p(z|x, \theta) - \log q(z|x, \lambda)] \\
 &= \arg \max_{\lambda} \underbrace{\mathbb{E}_{p(x, z|\theta)} [\log q(z|x, \lambda)]}_{\mathcal{R}(\lambda)} - \underbrace{\mathbb{E}_{p(x, z|\theta)} [\log p(z|x, \theta)]}_{\text{constant}}
 \end{aligned}$$

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\begin{aligned}
 & \arg \min_{\lambda} \mathbb{E}_{p(x)} [\text{KL} (p(z|x, \theta) \parallel q(z|x, \lambda))] \\
 &= \arg \min_{\lambda} \mathbb{E}_{p(x, z|\theta)} [\log p(z|x, \theta) - \log q(z|x, \lambda)] \\
 &= \arg \max_{\lambda} \underbrace{\mathbb{E}_{p(x, z|\theta)} [\log q(z|x, \lambda)]}_{\mathcal{R}(\lambda)} - \underbrace{\mathbb{E}_{p(x, z|\theta)} [\log p(z|x, \theta)]}_{\text{constant}}
 \end{aligned}$$

Gradient

$$\nabla_{\lambda} \mathcal{R}(\lambda) = \nabla_{\lambda} \mathbb{E}_{p(x, z|\theta)} [\log q(z|x, \lambda)]$$

Sleep Phase (Convenient) Objective

Flip the direction of the KL

$$\begin{aligned}
 & \arg \min_{\lambda} \mathbb{E}_{p(x)} [\text{KL} (p(z|x, \theta) \parallel q(z|x, \lambda))] \\
 &= \arg \min_{\lambda} \mathbb{E}_{p(x, z|\theta)} [\log p(z|x, \theta) - \log q(z|x, \lambda)] \\
 &= \arg \max_{\lambda} \underbrace{\mathbb{E}_{p(x, z|\theta)} [\log q(z|x, \lambda)]}_{\mathcal{R}(\lambda)} - \underbrace{\mathbb{E}_{p(x, z|\theta)} [\log p(z|x, \theta)]}_{\text{constant}}
 \end{aligned}$$

Gradient

$$\begin{aligned}
 \nabla_{\lambda} \mathcal{R}(\lambda) &= \nabla_{\lambda} \mathbb{E}_{p(x, z|\theta)} [\log q(z|x, \lambda)] \\
 &= \mathbb{E}_{p(x, z|\theta)} [\nabla_{\lambda} \log q(z|x, \lambda)]
 \end{aligned}$$

Sleep Phase (Convenient) Objective

Assumes fake data \tilde{x} and latent variables z to be fixed random draws from $p(x, z|\theta)$.

$$\arg \max_{\lambda} \mathbb{E}_{p(x, z|\theta)} [\log q(z|x, \lambda)]$$

Sleep Phase (Convenient) Objective

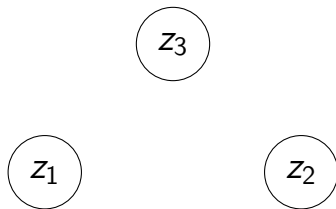
Assumes fake data \tilde{x} and latent variables z to be fixed random draws from $p(x, z|\theta)$.

$$\begin{aligned} & \arg \max_{\lambda} \mathbb{E}_{p(x, z|\theta)} [\log q(z|x, \lambda)] \\ & \approx^{\text{MC}} \arg \max_{\lambda} \log q(z|\tilde{x}, \lambda) \end{aligned}$$

where $z \sim p(z)$ and $\tilde{x} \sim p(x|z)$ (fake data!)

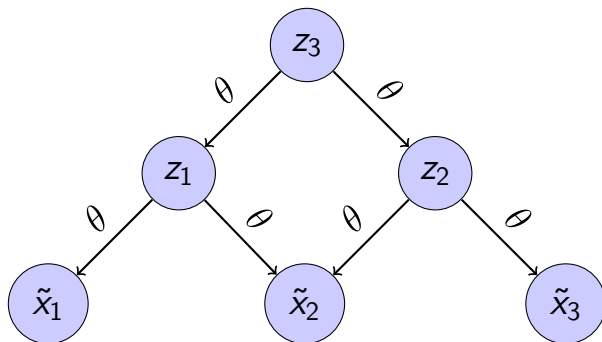
Sleep Phase Sampling

Sampling $(z, \tilde{x}) \sim p(x, z|\theta)$



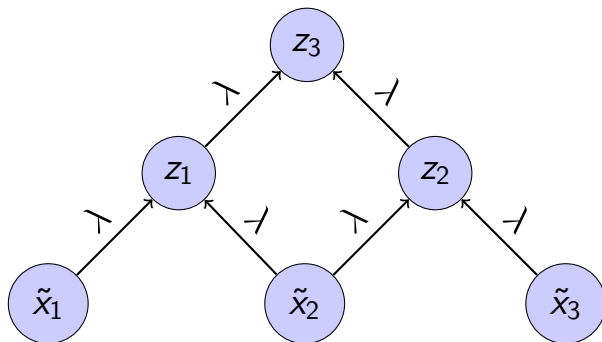
Sleep Phase Sampling

Sampling $(z, \tilde{x}) \sim p(x, z|\theta)$



Sleep Phase Update

Update λ



Wake-sleep Algorithm

Advantages

- ▶ Simple layer-wise updates
- ▶ Amortised inference: all latent variables are inferred from the same weights λ

Wake-sleep Algorithm

Advantages

- ▶ Simple layer-wise updates
- ▶ Amortised inference: all latent variables are inferred from the same weights λ

Drawbacks

- ▶ Inference and generative networks are trained on different objectives
- ▶ Inference weights λ are updated on fake data \tilde{x}
- ▶ Generative weights are bad initially, giving wrong signal to the updates of λ

First Attempt: Wake-Sleep

Neural Variational Inference

Score function estimator

Variance reduction

Variational Inference Learning (NVIL)

Generative model with NN likelihood

Variational Inference Learning (NVIL)

Generative model with NN likelihood

Let us consider a latent factor model for topic modelling:

Variational Inference Learning (NVIL)

Generative model with NN likelihood

Let us consider a latent factor model for topic modelling:

- ▶ a document $x = (x_1, \dots, x_N)$ consists of n i.i.d. categorical draws from that model

Variational Inference Learning (NVIL)

Generative model with NN likelihood

Let us consider a latent factor model for topic modelling:

- ▶ a document $x = (x_1, \dots, x_N)$ consists of n i.i.d. categorical draws from that model
- ▶ the categorical distribution in turn depends on binary latent factors $z = (z_1, \dots, z_K)$ which are also i.i.d.

Latent factor model

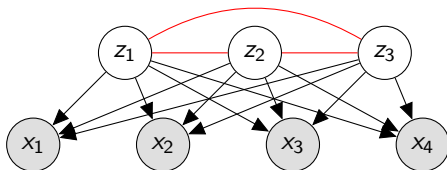
$$Z_j \sim \text{Bernoulli}(\phi) \quad (1 \leq k \leq K)$$

$$X_i|z \sim \text{Categorical}(f(z; \theta)) \quad (1 \leq i \leq N)$$

Here $0 < \phi < 1$ specifies a Bernoulli prior
and $f(\cdot; \theta)$ is a function computed by a
neural network with softmax output, e.g.

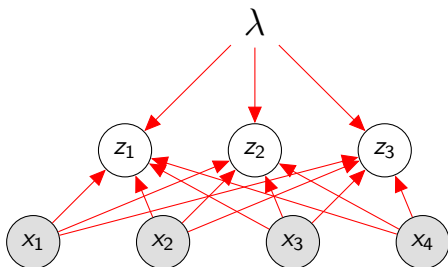
$$\begin{aligned} f(z; \theta) &= \text{softmax}(Wz + b) \\ \theta &= \{W, b\} \end{aligned}$$

Example Model



At inference time the latent variables are marginally dependent. For our variational distribution we are going to assume that they are not (recall: mean field assumption).

Mean Field Inference



The inference network needs to predict K Bernoulli parameters b_1^K . Any neural network with sigmoid output will do that job.

Inference Network

$$q(z|x, \lambda) = \prod_{k=1}^K \text{Bern}(z_k | b_k)$$

where $b_1^K = g(x; \lambda)$

Example architecture

$$h = \frac{1}{N} \sum_{i=1}^N E_{x_i} \quad b_1^K = \text{sigmoid}(Mh + c)$$

$$\lambda = \{E, M, c\}$$

Objective

$$\begin{aligned}\text{ELBO} &= \mathbb{E}_{q(z|x, \lambda)} [\log p(x, z|\theta)] + \mathbb{H}(q(z|x, \lambda)) \\ &= \mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] - \text{KL}(q(z|x, \lambda) \parallel p(z))\end{aligned}$$

Parameter estimation

$$\arg \max_{\theta, \lambda} \mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] - \text{KL}(q(z|x, \lambda) \parallel p(z))$$

KL

KL between K independent Bernoulli distributions is tractable

$$\begin{aligned}\text{KL} (q(z|x, \lambda) \parallel p(z|\phi)) &= \sum_{k=1}^K \text{KL} (q(z_k|x, \lambda) \parallel p(z_k|\phi)) \\ &= \sum_{k=1}^K b_k \log \frac{b_k}{\phi} + (1 - b_k) \log \frac{1 - b_k}{1 - \phi}\end{aligned}$$

Generative Network Gradient

$$\frac{\partial}{\partial \theta} \left(\mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] - \overbrace{\text{KL}(q(z|x, \lambda) || p(z))}^{\text{constant wrt } \theta} \right)$$

Generative Network Gradient

$$\begin{aligned}
 & \frac{\partial}{\partial \theta} \left(\mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] - \overbrace{\text{KL}(q(z|x, \lambda) || p(z))}^{\text{constant wrt } \theta} \right) \\
 &= \underbrace{\mathbb{E}_{q(z|x, \lambda)} \left[\frac{\partial}{\partial \theta} \log p(x|z, \theta) \right]}_{\text{expected gradient :)}}
 \end{aligned}$$

Generative Network Gradient

$$\frac{\partial}{\partial \theta} \left(\mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] - \overbrace{\text{KL}(q(z|x, \lambda) || p(z))}^{\text{constant wrt } \theta} \right)$$

$$= \underbrace{\mathbb{E}_{q(z|x, \lambda)} \left[\frac{\partial}{\partial \theta} \log p(x|z, \theta) \right]}_{\text{expected gradient :)}}$$

$$\stackrel{\text{MC}}{\approx} \frac{1}{S} \sum_{s=1}^S \frac{\partial}{\partial \theta} \log p(x|z^{(s)}, \theta) \quad \text{where } z^{(s)} \sim q(z|x, \lambda)$$

Inference Network Gradient

$$\frac{\partial}{\partial \lambda} \left(\mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] - \overbrace{\text{KL} (q(z|x, \lambda) || p(z))}^{\text{analytical}} \right)$$

Inference Network Gradient

$$\begin{aligned}
 & \frac{\partial}{\partial \lambda} \left(\mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] - \overbrace{\text{KL} (q(z|x, \lambda) || p(z))}^{\text{analytical}} \right) \\
 &= \frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] - \underbrace{\frac{\partial}{\partial \lambda} \text{KL} (q(z|x, \lambda) || p(z))}_{\text{analytical computation}}
 \end{aligned}$$

Inference Network Gradient

$$\begin{aligned}
 & \frac{\partial}{\partial \lambda} \left(\mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] - \overbrace{\text{KL} (q(z|x, \lambda) || p(z))}^{\text{analytical}} \right) \\
 &= \frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] - \underbrace{\frac{\partial}{\partial \lambda} \text{KL} (q(z|x, \lambda) || p(z))}_{\text{analytical computation}}
 \end{aligned}$$

The first term again requires approximation by sampling, but there is a problem

Inference Network Gradient

$$\frac{\partial}{\partial \lambda} \mathbb{E}_{q_{\lambda}(z|x)} [\log p_{\theta}(x|z)]$$

Inference Network Gradient

$$\begin{aligned} & \frac{\partial}{\partial \lambda} \mathbb{E}_{q_{\lambda}(z|x)} [\log p_{\theta}(x|z)] \\ &= \frac{\partial}{\partial \lambda} \sum_z q(z|x, \lambda) \log p(x|z, \theta) \end{aligned}$$

Inference Network Gradient

$$\begin{aligned}
 & \frac{\partial}{\partial \lambda} \mathbb{E}_{q_{\lambda}(z|x)} [\log p_{\theta}(x|z)] \\
 &= \frac{\partial}{\partial \lambda} \sum_z q(z|x, \lambda) \log p(x|z, \theta) \\
 &= \underbrace{\sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta)}_{\text{not an expectation}}
 \end{aligned}$$

Inference Network Gradient

$$\begin{aligned}
 & \frac{\partial}{\partial \lambda} \mathbb{E}_{q_{\lambda}(z|x)} [\log p_{\theta}(x|z)] \\
 &= \frac{\partial}{\partial \lambda} \sum_z q(z|x, \lambda) \log p(x|z, \theta) \\
 &= \underbrace{\sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta)}_{\text{not an expectation}}
 \end{aligned}$$

- MC estimator is non-differentiable

Inference Network Gradient

$$\begin{aligned}
 & \frac{\partial}{\partial \lambda} \mathbb{E}_{q_{\lambda}(z|x)} [\log p_{\theta}(x|z)] \\
 &= \frac{\partial}{\partial \lambda} \sum_z q(z|x, \lambda) \log p(x|z, \theta) \\
 &= \underbrace{\sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta)}_{\text{not an expectation}}
 \end{aligned}$$

- ▶ MC estimator is non-differentiable
- ▶ Differentiating the expression does not yield an expectation: cannot approximate via MC

Score function estimator

We can again use the log identity for derivatives

$$\begin{aligned} & \frac{\partial}{\partial \lambda} \mathbb{E}_{q_{\lambda}(z|x)} [\log p_{\theta}(x|z)] \\ &= \sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta) \end{aligned}$$

Score function estimator

We can again use the log identity for derivatives

$$\begin{aligned} & \frac{\partial}{\partial \lambda} \mathbb{E}_{q_{\lambda}(z|x)} [\log p_{\theta}(x|z)] \\ &= \sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta) \\ &= \sum_z q(z|x, \lambda) \frac{\partial}{\partial \lambda} (\log q(z|x, \lambda)) \log p(x|z, \theta) \end{aligned}$$

Score function estimator

We can again use the log identity for derivatives

$$\begin{aligned}
 & \frac{\partial}{\partial \lambda} \mathbb{E}_{q_{\lambda}(z|x)} [\log p_{\theta}(x|z)] \\
 &= \sum_z \frac{\partial}{\partial \lambda} (q(z|x, \lambda)) \log p(x|z, \theta) \\
 &= \sum_z q(z|x, \lambda) \frac{\partial}{\partial \lambda} (\log q(z|x, \lambda)) \log p(x|z, \theta) \\
 &= \underbrace{\mathbb{E}_{q(z|x, \lambda)} \left[\log p(x|z, \theta) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda) \right]}_{\text{expected gradient :)}}
 \end{aligned}$$

Score function estimator: remarks

We can now build an MC estimator

$$\begin{aligned} & \frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] \\ &= \mathbb{E}_{q(z|x, \lambda)} \left[\log p(x|z, \theta) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda) \right] \end{aligned}$$

Score function estimator: remarks

We can now build an MC estimator

$$\begin{aligned} & \frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] \\ &= \mathbb{E}_{q(z|x, \lambda)} \left[\log p(x|z, \theta) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda) \right] \\ &\stackrel{\text{MC}}{\approx} \frac{1}{S} \sum_{s=1}^S \log p(x|z^{(s)}, \theta) \frac{\partial}{\partial \lambda} \log q(z^{(s)}|x, \lambda) \end{aligned}$$

where $z^{(s)} \sim q(z|x, \lambda)$

Score function estimator: remarks

We can now build an MC estimator

$$\begin{aligned} & \frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] \\ &= \mathbb{E}_{q(z|x, \lambda)} \left[\log p(x|z, \theta) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda) \right] \\ &\stackrel{\text{MC}}{\approx} \frac{1}{S} \sum_{s=1}^S \log p(x|z^{(s)}, \theta) \frac{\partial}{\partial \lambda} \log q(z^{(s)}|x, \lambda) \\ &\text{where } z^{(s)} \sim q(z|x, \lambda) \end{aligned}$$

- magnitude of $\log p(x|z, \theta)$ varies widely

Score function estimator: remarks


We can now build an MC estimator

$$\begin{aligned}
 & \frac{\partial}{\partial \lambda} \mathbb{E}_{q(z|x, \lambda)} [\log p(x|z, \theta)] \\
 &= \mathbb{E}_{q(z|x, \lambda)} \left[\log p(x|z, \theta) \frac{\partial}{\partial \lambda} \log q(z|x, \lambda) \right] \\
 &\stackrel{\text{MC}}{\approx} \frac{1}{S} \sum_{s=1}^S \log p(x|z^{(s)}, \theta) \frac{\partial}{\partial \lambda} \log q(z^{(s)}|x, \lambda)
 \end{aligned}$$

where $z^{(s)} \sim q(z|x, \lambda)$

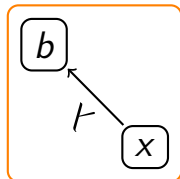
- magnitude of $\log p(x|z, \theta)$ varies widely
- model likelihood does not contribute to direction of gradient

Computation Graph



inference model

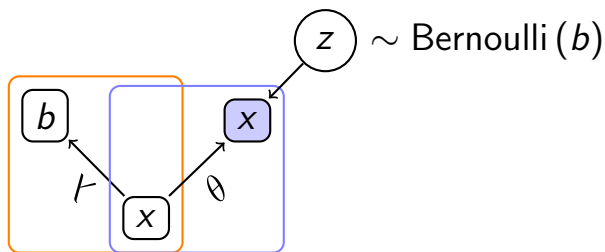
Computation Graph



$$z \sim \text{Bernoulli}(b)$$

inference model

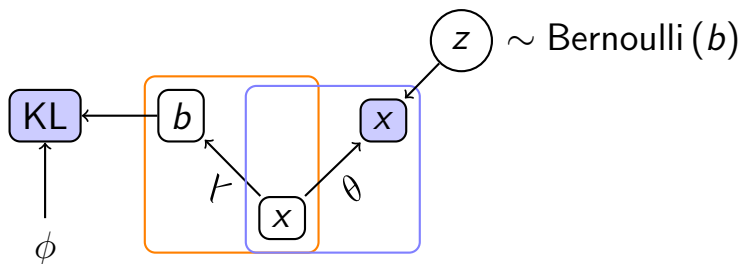
Computation Graph



inference model

generation model

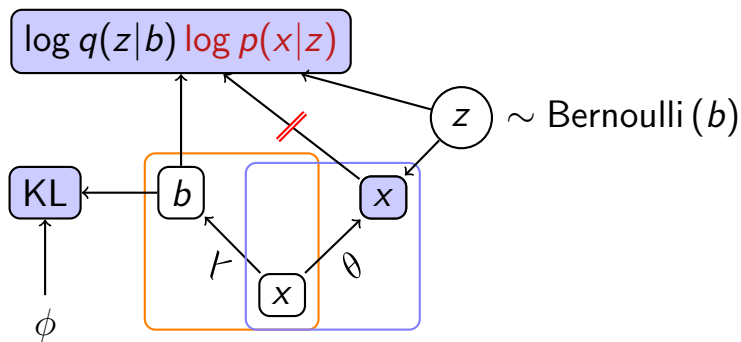
Computation Graph



inference model

generation model

Computation Graph



inference model

generation model

Pros and Cons

- ▶ Pros
 - ▶ Applicable to all distributions
 - ▶ Many libraries come with samplers for common distributions

Pros and Cons

- ▶ Pros
 - ▶ Applicable to all distributions
 - ▶ Many libraries come with samplers for common distributions
- ▶ Cons
 - ▶ High Variance!

First Attempt: Wake-Sleep

Neural Variational Inference

Score function estimator

Variance reduction

When variance is high we can

When variance is high we can

▶ sample more

When variance is high we can

- ▶ sample more
- ▶ use variance reduction techniques (e.g. baselines and control variates)

Control variates

Suppose we want to estimate $\mathbb{E}[f(Z)]$

$$\hat{f} \stackrel{\text{MC}}{\approx} \frac{1}{S} \sum_{s=1}^S f(z^{(s)})$$

and we know the expected value of another function $\psi(z)$ on the same support.

Control variates

Suppose we want to estimate $\mathbb{E}[f(Z)]$

$$\hat{f} \stackrel{\text{MC}}{\approx} \frac{1}{S} \sum_{s=1}^S f(z^{(s)})$$

and we know the expected value of another function $\psi(z)$ on the same support.

It holds that

$$\mathbb{E}[f(Z)] = \mathbb{E}[f(Z) - \psi(Z)] + \mathbb{E}[\psi(Z)]$$

Variance reduction

$$\hat{d} = \frac{1}{S} \left(\sum_{s=1}^S f(z^{(s)}) - \psi(z^{(s)}) \right) + \underbrace{\mathbb{E}[\psi(Z)]}_{\mu_\psi}$$

In general

$$\text{Var}(f - \psi) = \text{Var}(f) - 2 \text{Cov}(f, \psi) + \text{Var}(\psi)$$

If f and ψ are strongly correlated, then we improve on the original estimation problem:

$$\text{if } \text{Cov}(f, \psi) > \frac{\text{Var}(\psi)}{2}$$

Baselines

Fact

The Expectation of the score function is 0.

Baselines

Fact

The Expectation of the score function is 0.

$$\mathbb{E}_{q(z|x, \lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|x, \lambda) \right] = 0$$

Baselines

We attempt to centre the gradient estimate. To do this we learn a quantity C that we subtract from the reconstruction loss.

$$\mathbb{E}_{q(z|\lambda)} [\log q(z|\lambda) (\log p(x|z, \theta) - C)]$$

We call C a baseline. It does not change the expected gradient ([Williams, 1992](#)).

Baselines

$$\mathbb{E}_{q(z|\lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|\lambda) (\log p(x|z, \theta) - C) \right] =$$

Baselines

$$\mathbb{E}_{q(z|\lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|\lambda) (\log p(x|z, \theta) - C) \right] =$$

$$\underbrace{\mathbb{E}_{q(z|\lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|\lambda) \log p(x|z, \theta) \right]}_{\text{score function gradient}} -$$

Baselines

$$\begin{aligned}
 & \mathbb{E}_{q(z|\lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|\lambda) (\log p(x|z, \theta) - C) \right] = \\
 & \underbrace{\mathbb{E}_{q(z|\lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|\lambda) \log p(x|z, \theta) \right]}_{\text{score function gradient}} - \\
 & \underbrace{\mathbb{E}_{q(z|\lambda)} \left[\frac{\partial}{\partial \lambda} \log q(z|\lambda) \right]}_0 C
 \end{aligned}$$

Baselines

We can make baselines input-dependent to make them more flexible.

$$\log q(z|\lambda) (\log p(x|z, \theta) - C(x; \omega))$$

Baselines

We can make baselines input-dependent to make them more flexible.

$$\log q(z|\lambda) (\log p(x|z, \theta) - C(x; \omega))$$

However, baselines may not depend on the random value z ! Quantities that may depend on the random value ($C(z)$) are called **control variates**.

See [Blei et al. \(2012\)](#); [Ranganath et al. \(2014\)](#); [Gregor et al. \(2014\)](#).

Baselines

Baselines are predicted by a regression model (e.g. a neural net).

The model is trained using an L_2 -loss.

$$\min_{\omega} (C(x; \omega) - \log p(x|z, \theta))^2$$

Summary

- ▶ Wake-Sleep: train inference and generation networks with separate objectives

Summary

- ▶ Wake-Sleep: train inference and generation networks with separate objectives
- ▶ NVIL: a single objective (ELBO) for both models

Summary

- ▶ Wake-Sleep: train inference and generation networks with separate objectives
- ▶ NVIL: a single objective (ELBO) for both models
- ▶ Use score function estimator

Summary

- ▶ Wake-Sleep: train inference and generation networks with separate objectives
- ▶ NVIL: a single objective (ELBO) for both models
- ▶ Use score function estimator
- ▶ Always use baselines for variance reduction!

Implementation

Check one of our notebooks, e.g.

- ▶ inducing rationales for sentiment classification
[https://github.com/probabll/dgm4nlp/
tree/master/notebooks/sst](https://github.com/probabll/dgm4nlp/tree/master/notebooks/sst)

Literature I

David M. Blei, Michael I. Jordan, and John W. Paisley. Variational bayesian inference with stochastic search. In *ICML*, 2012. URL <http://icml.cc/2012/papers/687.pdf>.

Karol Gregor, Ivo Danihelka, Andriy Mnih, Charles Blundell, and Daan Wierstra. Deep autoregressive networks. In Eric P. Xing and Tony Jebara, editors, *ICML*, pages 1242–1250, 2014. URL <http://proceedings.mlr.press/v32/gregor14.html>.

Literature II

G. E. Hinton, P. Dayan, B. J. Frey, and R. M. Neal. The wake-sleep algorithm for unsupervised neural networks. *Science*, 268:1158–1161, 1995. URL <http://www.gatsby.ucl.ac.uk/~dayan/papers/hdfn95.pdf>.

Andriy Mnih and Karol Gregor. Neural variational inference and learning in belief networks. In *Proceedings of the 31st International Conference on International Conference on Machine Learning - Volume 32, ICML'14*, pages II–1791–II–1799.

Literature III

JMLR.org, 2014. URL <http://dl.acm.org/citation.cfm?id=3044805.3045092>.

Rajesh Ranganath, Sean Gerrish, and David Blei. Black Box Variational Inference. In Samuel Kaski and Jukka Corander, editors, *AISTATS*, pages 814–822, 2014. URL <http://proceedings.mlr.press/v33/ranganath14.pdf>.

Literature IV

Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3-4): 229–256, 1992. URL <https://doi.org/10.1007/BF00992696>.