Automatic Differentiation Variational Inference

Philip Schulz and Wilker Aziz
https://github.com/philschulz/VITutorial

What we know so far

- DGMs: probabilistic models parameterised by neural networks

What we know so far

- DGMs: probabilistic models parameterised by neural networks
- Objective: a lowerbound on likelihood

What we know so far

- DGMs: probabilistic models parameterised by neural networks
- Objective: a lowerbound on likelihood
- ELBO cannot be computed exactly, but we can obtain gradient estimates

What we know so far

- DGMs: probabilistic models parameterised by neural networks
- Objective: a lowerbound on likelihood
- ELBO cannot be computed exactly, but we can obtain gradient estimates
- Main problem: the MC estimator is not differentiable

What we know so far

- DGMs: probabilistic models parameterised by neural networks
- Objective: a lowerbound on likelihood
- ELBO cannot be computed exactly, but we can obtain gradient estimates
- Main problem: the MC estimator is not differentiable
    - Score function estimator: applicable to any model

What we know so far

- DGMs: probabilistic models parameterised by neural networks
- Objective: a lowerbound on likelihood
- ELBO cannot be computed exactly, but we can obtain gradient estimates
- Main problem: the MC estimator is not differentiable
  - Score function estimator: applicable to any model
  - Reparameterised gradients: so far seems applicable only to Gaussian variables

Reparameterised gradients: Gaussian

We have seen one case, namely,
if $\epsilon \sim \mathcal{N}(0, I)$ and $Z \sim \mathcal{N}(\mu, \sigma^2)$

Reparameterised gradients: Gaussian

We have seen one case, namely,
if $\epsilon \sim \mathcal{N}(0, I)$ and $Z \sim \mathcal{N}(\mu, \sigma^2)$
Then

$$Z \sim \mu + \sigma\epsilon$$

and

$$\frac{\partial}{\partial\lambda}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right]$$

Reparameterised gradients: Gaussian

We have seen one case, namely,
if $\epsilon \sim \mathcal{N}(0, I)$ and $Z \sim \mathcal{N}(\mu, \sigma^2)$
Then

$$Z \sim \mu + \sigma\epsilon$$

and

$$\frac{\partial}{\partial\lambda}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right]$$
$$= \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial\lambda}g(z = \mu + \sigma\epsilon)\right]$$

Reparameterised gradients: Gaussian

We have seen one case, namely,
if $\epsilon \sim \mathcal{N}(0, I)$ and $Z \sim \mathcal{N}(\mu, \sigma^2)$
Then

$$Z \sim \mu + \sigma\epsilon$$

and

$$\begin{aligned}
&\frac{\partial}{\partial\lambda}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right] \\
&= \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial\lambda}g(z = \mu + \sigma\epsilon)\right] \\
&= \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z = \mu + \sigma\epsilon)\frac{\partial z}{\partial\lambda}\right]
\end{aligned}$$

Reparameterised gradients: Gaussian

Location

$$\frac{\partial}{\partial \mu}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right] = \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z = \mu + \sigma\epsilon)\frac{\partial z}{\partial \mu}\right]$$

Reparameterised gradients: Gaussian

Location

$$\frac{\partial}{\partial \mu}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right] = \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\frac{\partial z}{\partial \mu}\right]$$
$$= \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\right]$$

Reparameterised gradients: Gaussian

Location

$$\frac{\partial}{\partial \mu}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right] = \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\frac{\partial z}{\partial \mu}\right]$$
$$= \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\right]$$

Scale

$$\frac{\partial}{\partial \sigma}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right] = \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\frac{\partial z}{\partial \sigma}\right]$$

Reparameterised gradients: Gaussian

Location

$$\frac{\partial}{\partial \mu}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right] = \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\frac{\partial z}{\partial \mu}\right]$$

$$= \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\right]$$

Scale

$$\frac{\partial}{\partial \sigma}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right] = \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\frac{\partial z}{\partial \sigma}\right]$$

$$= \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\epsilon\right]$$

Reparameterised gradients: Gaussian

Location

$$\frac{\partial}{\partial\mu}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right] = \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\frac{\partial z}{\partial\mu}\right]$$

$$= \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\right]$$

Scale

$$\frac{\partial}{\partial\sigma}\mathbb{E}_{\mathcal{N}(z|\mu,\sigma^2)}\left[g(z)\right] = \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\frac{\partial z}{\partial\sigma}\right]$$

$$= \mathbb{E}_{\mathcal{N}(0,I)}\left[\frac{\partial}{\partial z}g(z=\mu+\sigma\epsilon)\epsilon\right]$$

Multivariate calculus recap

Let $x \in \mathbb{R}^D$ and let $\mathcal{T} : \mathbb{R}^D \to \mathbb{R}^D$ be differentiable and invertible

- $y = \mathcal{T}(x)$
- $x = \mathcal{T}^{-1}(y)$

Jacobian

The Jacobian matrix $\mathbf{J} = J_{\mathcal{T}}(x)$ of $\mathcal{T}$ assessed at x is the matrix of partial derivatives

$$J_{ij} = \frac{\partial y_i}{\partial x_j}$$

Jacobian

The Jacobian matrix $\mathbf{J} = J_{\mathcal{T}}(x)$ of $\mathcal{T}$ assessed at x is the matrix of partial derivatives

$$J_{ij} = \frac{\partial y_i}{\partial x_j}$$

Inverse function theorem

$$J_{\mathcal{T}^{-1}}(y) = \left(J_{\mathcal{T}}(x)\right)^{-1}$$

Integration by substitution

We can integrate a function $g(x)$
by substituting $x = \mathcal{T}^{-1}(y)$

$$\int g(x)\mathrm{d}x$$

Integration by substitution

We can integrate a function $g(x)$
by substituting $x = \mathcal{T}^{-1}(y)$

$$\int g(x)\mathrm{d}x = \int g(\mathcal{T}^{-1}(y)) \left|\det J_{\mathcal{T}^{-1}}(y)\right| \mathrm{d}y$$

Integration by substitution

We can integrate a function $g(x)$
by substituting $x = \mathcal{T}^{-1}(y)$

$$\int g(x)\mathrm{d}x = \int g(\mathcal{T}^{-1}(y)) \left|\det J_{\mathcal{T}^{-1}}(y)\right| \mathrm{d}y$$

and similarly for a function $h(y)$

$$\int h(y)\mathrm{d}y$$

Integration by substitution

We can integrate a function $g(x)$ by substituting $x = \mathcal{T}^{-1}(y)$

$$\int g(x)\mathrm{d}x = \int g(\mathcal{T}^{-1}(y)) \left|\det J_{\mathcal{T}^{-1}}(y)\right| \mathrm{d}y$$

and similarly for a function $h(y)$

$$\int h(y)\mathrm{d}y = \int h(\mathcal{T}(x)) \left|\det J_{\mathcal{T}}(x)\right| \mathrm{d}x$$

Integration by substitution

We can integrate a function $g(x)$
by substituting $x = \mathcal{T}^{-1}(y)$

$$\int g(x)\mathrm{d}x = \int g(\mathcal{T}^{-1}(y)) \left|\det J_{\mathcal{T}^{-1}}(y)\right| \mathrm{d}y$$

and similarly for a function $h(y)$

$$\int h(y)\mathrm{d}y = \int h(\mathcal{T}(x)) \left|\det J_{\mathcal{T}}(x)\right| \mathrm{d}x$$

Rule of thumb

$$\mathrm{d}x = \left|\det J_{\mathcal{T}^{-1}}(y)\right| \mathrm{d}y$$

Integration by substitution

We can integrate a function $g(x)$
by substituting $x = \mathcal{T}^{-1}(y)$

$$\int g(x)\mathrm{d}x = \int g(\mathcal{T}^{-1}(y)) \left|\det J_{\mathcal{T}^{-1}}(y)\right| \mathrm{d}y$$

and similarly for a function $h(y)$

$$\int h(y)\mathrm{d}y = \int h(\mathcal{T}(x)) \left|\det J_{\mathcal{T}}(x)\right| \mathrm{d}x$$

Rule of thumb

$$\mathrm{d}x = \left|\det J_{\mathcal{T}^{-1}}(y)\right| \mathrm{d}y \qquad \mathrm{d}y = \left|\det J_{\mathcal{T}}(x)\right| \mathrm{d}x$$

Change of density

Let X take on values in $\mathbb{R}^D$ with density $f_X(x)$

Change of density

Let X take on values in $\mathbb{R}^D$ with density $f_X(x)$ and recall that $y = \mathcal{T}(x)$ and $x = \mathcal{T}^{-1}(y)$

Change of density

Let X take on values in $\mathbb{R}^D$ with density $f_X(x)$ and recall that $y = \mathcal{T}(x)$ and $x = \mathcal{T}^{-1}(y)$

Then $\mathcal{T}$ induces a density $f_Y(y)$ expressed as

$$f_Y(y) = f_X(x = \mathcal{T}^{-1}(y)) \left|\det J_{\mathcal{T}^{-1}}(y)\right|$$

Change of density

Let X take on values in $\mathbb{R}^D$ with density $f_X(x)$ and recall that $y = \mathcal{T}(x)$ and $x = \mathcal{T}^{-1}(y)$

Then $\mathcal{T}$ induces a density $f_Y(y)$ expressed as

$$f_Y(y) = f_X(x = \mathcal{T}^{-1}(y)) \left|\det J_{\mathcal{T}^{-1}}(y)\right|$$

and then it follows that

$$f_X(x) = f_Y(y = \mathcal{T}(x)) \left|\det J_{\mathcal{T}}(x)\right|$$

Revisting reparameterised gradients

Let Z take on values in $\mathbb{R}^D$ with pdf $q(z|\lambda)$

Revisting reparameterised gradients

Let Z take on values in $\mathbb{R}^D$ with pdf $q(z|\lambda)$

The idea is to count on a *standardisation* procedure

Revisting reparameterised gradients

Let Z take on values in $\mathbb{R}^D$ with pdf $q(z|\lambda)$

The idea is to count on a *standardisation* procedure
a transformation $\mathcal{S}_\lambda : \mathbb{R}^D \to \mathbb{R}^D$ such that

Revisting reparameterised gradients

Let Z take on values in $\mathbb{R}^D$ with pdf $q(z|\lambda)$

The idea is to count on a *standardisation* procedure
a transformation $\mathcal{S}_\lambda : \mathbb{R}^D \to \mathbb{R}^D$ such that

$$\mathcal{S}_\lambda(z) \sim \pi(\epsilon)$$
$$\mathcal{S}_\lambda^{-1}(\epsilon) \sim f(z|\lambda)$$

- $\pi(\epsilon)$ does not depend on parameters λ
  we call it a *standard* density

Revisting reparameterised gradients

Let Z take on values in $\mathbb{R}^D$ with pdf $q(z|\lambda)$

The idea is to count on a *standardisation* procedure
a transformation $\mathcal{S}_\lambda : \mathbb{R}^D \to \mathbb{R}^D$ such that

$$\mathcal{S}_\lambda(z) \sim \pi(\epsilon)$$
$$\mathcal{S}_\lambda^{-1}(\epsilon) \sim f(z|\lambda)$$

- $\pi(\epsilon)$ does not depend on parameters λ
  we call it a *standard* density
- $\mathcal{S}_\lambda(z)$ absorbs dependency on λ

Reparameterised expectations

If we are interested in

$$\mathbb{E}_{q(z|\lambda)}\left[g(z)\right]$$

Reparameterised expectations

If we are interested in

$$\mathbb{E}_{q(z|\lambda)}\left[g(z)\right] = \int q(z|\lambda)g(z)\mathrm{d}z$$

Reparameterised expectations

If we are interested in

$$\mathbb{E}_{q(z|\lambda)}\left[g(z)\right] = \int q(z|\lambda)g(z)\mathrm{d}z$$

$$= \int \underbrace{\pi(\mathcal{S}_\lambda(z))\left|\det J_{\mathcal{S}_\lambda}(z)\right|}_{\text{change of density}} g(z)\mathrm{d}z$$

Reparameterised expectations

If we are interested in

$$\mathbb{E}_{q(z|\lambda)}\left[g(z)\right] = \int q(z|\lambda) g(z) \mathrm{d}z$$

$$= \int \underbrace{\pi(\mathcal{S}_\lambda(z)) \left|\det J_{\mathcal{S}_\lambda}(z)\right|}_{\text{change of density}} g(z) \mathrm{d}z$$

$$= \int \pi(\epsilon) \left|\det J_{\mathcal{S}_\lambda^{-1}}(\epsilon)\right|^{-1}$$

Reparameterised expectations

If we are interested in

$$\mathbb{E}_{q(z|\lambda)}\left[g(z)\right] = \int q(z|\lambda) g(z) \mathrm{d}z$$

$$= \int \underbrace{\pi(\mathcal{S}_\lambda(z)) \left|\det J_{\mathcal{S}_\lambda}(z)\right|}_{\text{change of density}} g(z) \mathrm{d}z$$

$$= \int \pi(\epsilon) \left|\det J_{\mathcal{S}_\lambda^{-1}}(\epsilon)\right|^{-1} g(\mathcal{S}_\lambda^{-1}(\epsilon))$$

Reparameterised expectations

If we are interested in

$$\mathbb{E}_{q(z|\lambda)}[g(z)] = \int q(z|\lambda) g(z) \mathrm{d}z$$

$$= \int \underbrace{\pi(\mathcal{S}_\lambda(z)) \left|\det J_{\mathcal{S}_\lambda}(z)\right|}_{\text{change of density}} g(z) \mathrm{d}z$$

$$= \int \pi(\epsilon) \left|\det J_{\mathcal{S}_\lambda^{-1}}(\epsilon)\right|^{-1} g(\mathcal{S}_\lambda^{-1}(\epsilon)) \left|\det J_{\mathcal{S}_\lambda^{-1}}(\epsilon)\right| \mathrm{d}\epsilon$$

Reparameterised expectations

If we are interested in

$$\mathbb{E}_{q(z|\lambda)}[g(z)] = \int q(z|\lambda) g(z) \mathrm{d}z$$
$$= \int \underbrace{\pi(\mathcal{S}_\lambda(z)) \left|\det J_{\mathcal{S}_\lambda}(z)\right|}_{\text{change of density}} g(z) \mathrm{d}z$$
$$= \int \pi(\epsilon) \left|\det J_{\mathcal{S}_\lambda^{-1}}(\epsilon)\right|^{-1} g(\mathcal{S}_\lambda^{-1}(\epsilon)) \left|\det J_{\mathcal{S}_\lambda^{-1}}(\epsilon)\right| \mathrm{d}\epsilon$$
$$= \int \pi(\epsilon) g(\mathcal{S}_\lambda^{-1}(\epsilon)) \mathrm{d}\epsilon$$

Reparameterised expectations

If we are interested in

$$\mathbb{E}_{q(z|\lambda)}\left[g(z)\right] = \int q(z|\lambda) g(z) \mathrm{d}z$$

$$= \int \underbrace{\pi(\mathcal{S}_\lambda(z)) \left|\det J_{\mathcal{S}_\lambda}(z)\right|}_{\text{change of density}} g(z) \mathrm{d}z$$

$$= \int \pi(\epsilon) \left|\det J_{\mathcal{S}_\lambda^{-1}}(\epsilon)\right|^{-1} g(\mathcal{S}_\lambda^{-1}(\epsilon)) \left|\det J_{\mathcal{S}_\lambda^{-1}}(\epsilon)\right| \mathrm{d}\epsilon$$

$$= \int \pi(\epsilon) g(\mathcal{S}_\lambda^{-1}(\epsilon)) \mathrm{d}\epsilon = \mathbb{E}_{\pi(\epsilon)}\left[g(\mathcal{S}_\lambda^{-1}(\epsilon))\right]$$

Reparameterised gradients

For optimisation, we need tractable gradients

$$\frac{\partial}{\partial \lambda}\mathbb{E}_{q(z|\lambda)}\left[g(z)\right] = \frac{\partial}{\partial \lambda}\mathbb{E}_{\pi(\epsilon)}\left[g(\mathcal{S}_{\lambda}^{-1}(\epsilon))\right]$$

Reparameterised gradients

For optimisation, we need tractable gradients

$$\frac{\partial}{\partial \lambda}\mathbb{E}_{q(z|\lambda)}\left[g(z)\right] = \frac{\partial}{\partial \lambda}\mathbb{E}_{\pi(\epsilon)}\left[g(\mathcal{S}_{\lambda}^{-1}(\epsilon))\right]$$

since now the measure of integration does not depend on λ, we can obtain a gradient estimate

$$\frac{\partial}{\partial \lambda}\mathbb{E}_{q(z|\lambda)}\left[g(z)\right] = \mathbb{E}_{\pi(\epsilon)}\left[\frac{\partial}{\partial \lambda}g(\mathcal{S}_{\lambda}^{-1}(\epsilon))\right]$$

Reparameterised gradients

For optimisation, we need tractable gradients

$$\frac{\partial}{\partial \lambda}\mathbb{E}_{q(z|\lambda)}\left[g(z)\right] = \frac{\partial}{\partial \lambda}\mathbb{E}_{\pi(\epsilon)}\left[g(\mathcal{S}_{\lambda}^{-1}(\epsilon))\right]$$

since now the measure of integration does not depend on λ, we can obtain a gradient estimate

$$\frac{\partial}{\partial \lambda}\mathbb{E}_{q(z|\lambda)}\left[g(z)\right] = \mathbb{E}_{\pi(\epsilon)}\left[\frac{\partial}{\partial \lambda}g(\mathcal{S}_{\lambda}^{-1}(\epsilon))\right]$$

$$\overset{\text{MC}}{\approx} \frac{1}{M}\sum_{\substack{i=1 \\ \epsilon_i \sim \pi(\epsilon)}}^{M} \frac{\partial}{\partial \lambda}g(\mathcal{S}_{\lambda}^{-1}(\epsilon_i))$$

Standardisation functions

Location-scale family

- a family of distributions where for
  $F_X(x) = \Pr\{X \leq x\}$
  if $Y = a + bX$, then $F_Y(y|a, b) = F_X(\frac{z-a}{b})$

Standardisation functions

Location-scale family

- a family of distributions where for $F_X(x) = \Pr\{X \leq x\}$
  if $Y = a + bX$, then $F_Y(y|a, b) = F_X(\frac{z-a}{b})$
- if we can draw from $f_X(x)$, we can draw from $f_Y(y|a, b)$

Standardisation functions

Location-scale family

- a family of distributions where for
  $F_X(x) = \Pr\{X \leq x\}$
  if $Y = a + bX$, then $F_Y(y|a, b) = F_X(\frac{z-a}{b})$
- if we can draw from $f_X(x)$, we can draw from $f_Y(y|a, b)$
- the transformation absorbs the parameters a, b

Standardisation functions

Location-scale family

- a family of distributions where for
  $F_X(x) = \Pr\{X \leq x\}$
  if $Y = a + bX$, then $F_Y(y|a,b) = F_X(\frac{z-a}{b})$
- if we can draw from $f_X(x)$, we can draw from $f_Y(y|a,b)$
- the transformation absorbs the parameters a, b

Examples: Gaussian, Laplace, Cauchy, Uniform

Standardisation functions (cont.)

Inverse cdf

- for univariate Z with pdf $f_Z(z)$ and cdf $F_Z(z)$

$$P \sim \mathcal{U}(0,1) \qquad Z \sim F_Z^{-1}(P)$$

where $F_Z^{-1}(p)$ is the *quantile function*

Standardisation functions (cont.)

Inverse cdf

- for univariate Z with pdf $f_Z(z)$ and cdf $F_Z(z)$

$$P \sim \mathcal{U}(0,1) \qquad Z \sim F_Z^{-1}(P)$$

where $F_Z^{-1}(p)$ is the *quantile function*

Gumbel distribution

- $f_Z(z|\mu,\beta) = \beta^{-1}\exp(-z - \exp(-z))$
- $F_Z(z|\mu,\beta) = \exp\left(-\exp\left(-\frac{z-\mu}{\beta}\right)\right)$
- $F_Z^{-1}(p) = \mu - \beta\log(-\log p)$

Beyond

Many interesting densities are not location-scale families

- e.g. Beta, Gamma

The inverse cdf of a multivariate rv is seldom known in closed-form

- Dirichlet, von Mises-Fisher

Automatic Differentiation VI

Motivation

- many models have intractable posteriors
  their normalising constants (evidence) lacks
  analytic solutions

Automatic Differentiation VI

Motivation

- many models have intractable posteriors
  their normalising constants (evidence) lacks analytic solutions
- but many models are differentiable
  that's the main constraint for using NNs

Automatic Differentiation VI

Motivation

- many models have intractable posteriors
  their normalising constants (evidence) lacks analytic solutions
- but many models are differentiable
  that's the main constraint for using NNs

Reparameterised gradients are a step towards automatising VI for differentiable models

Automatic Differentiation VI

Motivation

- many models have intractable posteriors
  their normalising constants (evidence) lacks analytic solutions
- but many models are differentiable
  that's the main constraint for using NNs

Reparameterised gradients are a step towards automatising VI for differentiable models

- but not every model of interest employs rvs for which a standardisation function is known

Example

Suppose we have some ordinal data which we assume to be Poisson-distributed

$$X|\lambda \sim \text{Poisson}(\lambda)$$

and suppose we want to impose

Differentiable models

We focus on *differentiable probability models*

$$p(x, z) = p(x|z)p(z)$$

Differentiable models

We focus on *differentiable probability models*

$$p(x, z) = p(x|z)p(z)$$

- members of this class have continuous latent variables z

Differentiable models

We focus on *differentiable probability models*

$$p(x, z) = p(x|z)p(z)$$

- members of this class have continuous latent variables z
- and the gradient $\nabla_z \log p(x, z)$ is valid within the *support* of the prior $\mathrm{supp}(p(z)) = \{z \in \mathbb{R}^K : p(z) > 0\} \subseteq \mathbb{R}^K$

VI optimisation problem

Let's focus on the design and optimisation of the variational approximation

$$\arg\min_{q(z)} \mathrm{KL}\left(q(z) \;||\; p(z|x)\right)$$

VI optimisation problem

Let's focus on the design and optimisation of the variational approximation

$$\underset{q(z)}{\arg\min}\, \mathrm{KL}\left(q(z) \;||\; p(z|x)\right)$$

Support matching constraint

To automatise the search for a variational approximation $q(z)$ we must ensure that

$$\mathrm{supp}(q(z)) \subseteq \mathrm{supp}(p(z|x))$$

- otherwise KL is not defined

Support matching constraint

To automatise the search for a variational approximation $q(z)$ we must ensure that

$$\mathrm{supp}(q(z)) \subseteq \mathrm{supp}(p(z|x))$$

- otherwise KL is not defined
- but what is the support of $p(z|x)$?

Support matching constraint

To automatise the search for a variational approximation $q(z)$ we must ensure that

$$\mathrm{supp}(q(z)) \subseteq \mathrm{supp}(p(z|x))$$

- otherwise KL is not defined
- but what is the support of $p(z|x)$?
- typically the same as the support of $p(z)$ as long as $p(x, z) > 0$ if $p(z) > 0$

Parametric family

So let's constrain $q(z)$ to a family $\mathcal{Q}$ whose support is included in the support of the prior

$$\mathcal{Q} = \{q(z;\phi) : \phi \in \Phi\}$$

- a parameter vector ϕ picks out a member $q(z;\phi)$ of the family

Our objective now is

$$\arg\min_{\phi\in\Phi} \mathrm{KL}\left(q(z;\phi) \;||\; p(z|x)\right)$$

ADVI

Minimising the original VI problem is equivalent to maximising the ELBO

$$\mathcal{L}(\phi) = \mathbb{E}_{q(z;\phi)}\left[\log p(x, z)\right] + \mathbb{H}\left(q(z|\phi)\right) \qquad (1)$$

From the point of view of a black-box procedure, this objective poses two problems

1. intractable expectations

ADVI

Minimising the original VI problem is equivalent to maximising the ELBO

$$\mathcal{L}(\phi) = \mathbb{E}_{q(z;\phi)}\left[\log p(x, z)\right] + \mathbb{H}\left(q(z|\phi)\right) \qquad (1)$$

From the point of view of a black-box procedure, this objective poses two problems

1. intractable expectations Reparameterised Gradients!
2. custom $\mathrm{supp}(q(z;\phi))$

Idea

1. let's find a way to transform $\mathrm{supp}(p(z))$ to the complete real coordinate space